

ECE 8870 Project 2

Qazi Zarif Ul Islam

Pawprint: qzic2d

University of Missouri-columbia

04/21/2024

1 Technical Description

In this project, an RNN and an LSTM were implemented using pytorch. The broader goal of this project was to understand the working principles of RNNs and LSTMs and gain a comparative understanding of the two neural network architectures.

In the following sections, first the basic components of an RNN and an LSTM are introduced as well as an explanation of how backpropagation occurs in an RNN. This explanation will be more qualitative than mathematically rigorous. A more extensive treatment of backpropagation in RNNs can be found in —. After that, the experiments and results shall be demonstrated and discussed before a final conclusion section that talks about what more can be done to understand RNNs and LSTMs and the deficiencies of this project.

The code is available at : https://github.com/Murdock135/neural-nets-at-mizzou/tree/main/ECE_8770/project_2

1.1 Recurrent Neural Networks

Recurrent Neural Networks are neural neural networks that store the hidden layer’s output at the current time step so that it can influence the output of the hidden layer at the next time step. Qualitatively, this is commonly thought of as information being passed to the next time step. Figure 1 shows this in two forms. On the left the rnn is said to be in “rolled”(in time) form whereas on the right is is said to be in “unrolled”(in time) form. The figure tells us that the rnn can be seen as a “feedback” network of sorts where the output of the hidden layer h_t is passed back to the hidden layer. The left picture fails to capture that h_t is passed to the next time step. The unrolled version captures this well. On the other hand, the rolled version captures a very crucial fact; that throughout the hidden layers at different time steps, *there is only one weight matrix*, shared throughout the entire RNN at all time steps. Figure ?? explicitly illustrates this.

The shaded blue boxes are each a Multi-layer perceptron. The only caveat here is that in addition to the inputs of the “current time step” x_t , it also recieves *as* input, the weighted outputs of the previous hidden layer h_{t-1} . Thus, the input to the current hidden layer is expressed as a concatenation $[h_{t-1}; x_t]$. Thus the equations defining the RNN are,

$$h_t = \phi(W_{hh}h_{t-1} + W_{xh}x_t + b_h) \tag{1}$$




Figure 1: [2]: On the left recurrent connections are depicted via cyclic edges. On the right, we unfold the RNN over time steps. Here, recurrent edges span adjacent time steps, while conventional connections are computed synchronously

$$o_t = W_o h_t + b_o \quad (2)$$

Where in equation 1, W_{hh} indicates **the weights going from the hidden layer at the previous time step to the hidden layer at the current time step**, W_{xh} indicates **the weights going from the current inputs to the current hidden layer**. In equation 2 W_o indicates **the weights going from the hidden layer at the current time step to the output layer of the current time step**.

With these equations, the learning rule the RNN can be derived for any optimization algorithm e.g. stochastic gradient descent, adaptive momentum, RMS prop, etc. The algorithms won't be discussed in this report but the gradients of the error wrt the learnable parameters are mentioned in the next section, **backpropagation through time**.

1.1.1 Backpropagation through time (Backpropagation for RNNs)

For an RNN, the total loss is expressed as,

$$\hat{L} = \frac{1}{T} \sum_{t=1}^T \ell(y_t, d_t) = \frac{1}{T} (l_1 + l_2 + \dots + l_t \dots + l_T) \quad (3)$$

Where the small l_t 's indicate loss at a particular time step. The gradients of this total loss with respect to the learnable parameters are,

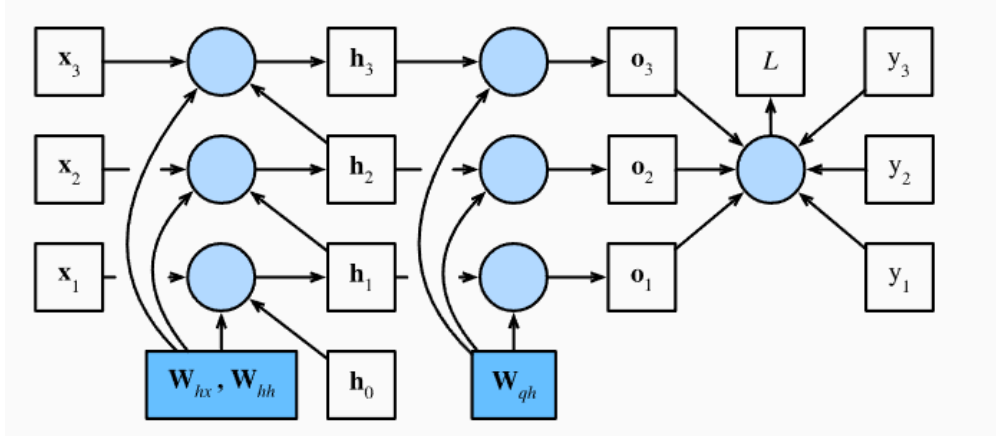


Figure 2: [2]: Computational graph showing dependencies for an RNN model with three time steps. Boxes represent variables (not shaded) or parameters (shaded) and circles represent operators.

$$\frac{\partial l_t}{\partial w_h} = \frac{\partial l_t}{\partial o_t} \frac{\partial o_t}{\partial h_t} \frac{\partial h_t}{\partial w_h} \quad (4)$$

$$\frac{\partial h_t}{\partial w_h} = \frac{\partial h_t}{\partial w_h} + \sum_{i=1}^{t-1} \prod_{j-1=i}^t \left(\frac{\partial h_j}{\partial h_{j-1}} \right) \frac{\partial h_i}{\partial w_h} \quad (5)$$

$$\frac{\partial L}{\partial W_{oh}} = \sum_t^T \frac{\partial L}{\partial o_t} h_t \quad (6)$$

$$\frac{\partial L}{\partial W_{hx}} = \sum_t^T \frac{\partial L}{\partial h_t} x_t \quad (7)$$

$$\frac{\partial L}{\partial w_h} = \sum_t^T \frac{\partial l_t}{\partial o_t} \frac{\partial o_t}{\partial h_t} \frac{\partial h_t}{\partial w_h} \quad (8)$$

Where $w_h = [W_{hx}; W_{hh}]$ is the concatenated weight matrix, W_{oh} is the weight matrix from the hidden layer to the output layer, W_{hx} is the weight matrix from the input layer to the hidden layer.

Observe 5. Any gradient based algorithm would compute this gradient and when it does, it has to recursively calculate the change of h_j wrt h_{j-1} and if the sequence is long, this product term will either be very large (if each gradient term > 1) or be miniscule (if each gradient term < 1). These two situations are called the *exploding gradient* and the *diminishing or vanishing gradient* problems, respectively. To solve this issue, several techniques are used to truncate this product. [2] mentions two such techniques; (1) *Regular truncation* and (2) *Randomized truncation*. The equations won't be derived in this report but an illustration of the three

techniques is presented in fig 3. In regular truncated BPTT, the sum is truncated after τ steps whereas in randomized BPTT, the initial sequences are split up into sub-sequences of “random” length (here, the length of the subsequences are cast as random variables). [1]




Figure 3: Comparing strategies for computing gradients in RNNs. From top to bottom: randomized truncation, regular truncation, and full computation.

1.2 Long Short Term Memory nets (LSTMs)

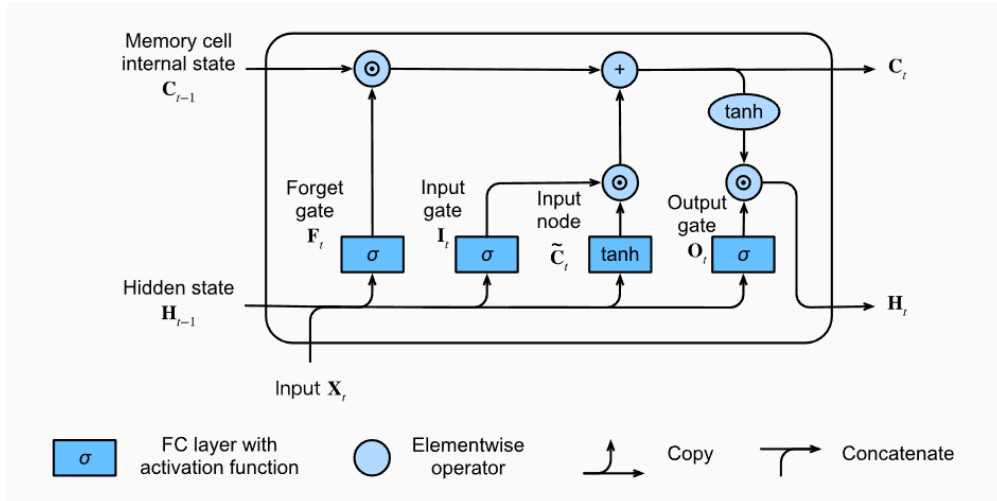


Figure 4: lstms

LSTMs

1.3 RNN/LSTM configurations

2 Experiments and Results

2.1 Training details

The following experiments were run on an NVIDIA CUDA RTX 3050 laptop GPU. The machine has a total memory of 16 GigaBytes. Each experiment was run for 100 epochs.

References

- [1] Corentin Tallec and Yann Ollivier. *Unbiasing Truncated Backpropagation Through Time*. 2017. arXiv: 1705.08209 [cs.NE].
- [2] Aston Zhang et al. *Dive into Deep Learning*. https://d2l.ai/chapter_recurrent-neural-networks/index.html. Cambridge University Press, 2023. Chap. 9.